

Processor Documentation

Ruben Saunders

August 2024

1 Principals

- This processor will operate a RISC instruction set.
- This processor has a word size of 64 bits, and supports both floats (4 bytes) and doubles (8 bytes).
- The instruction set will provide methods to load values into and out of registers. Then, most operations will be on registers.
- Load/store instructions operate on 32-bit immediates.
- Arithmetic and logic instructions operate on full registers, so 64-bit.

2 Registers

See below for a list of registers. All registers are 64-bit. Register names are preceded by a dollar ‘\$’ sign.

Symbol	Name	Bit	Description
Special Registers			
\$ip	Instruction Pointer		Point to next address to execute as an instruction.
\$sp	Stack Pointer		Top address of the stack.
\$fp	Frame Pointer		Point to the next byte beyond the last stack frame.
\$flag	Flag Register	6-64	
		5	Error status: 1=error, 0=ok. Indicates if exited gracefully.
		4	Execution status: 1=executing, 0=halted.
		3	Zero flag. Indicates if register is zero. Updated on most instructions’ dest register.
		0-2	Comparison bits. <ul style="list-style-type: none">• 000: not equal.• 001: equal.• 010: less than.• 011: less than or equal to.• 110: greater than.• 111: greater than or equal to.
\$ret	Return Value Register		Contains value returned from function, syscall, etc. Contains process exit code on halt.
\$zero	Zero		Hardwired to zero.
General Purpose Registers			
\$r1 – \$r16	GPRs		Register for general use.
\$s1 – \$s8	Preserved GPRs		Register for general use. Values are preserved in stack frame.

3 Addressing Modes

An argument can have the following type.

Indicator	Name	Syntax	Operation	Size
00	Immediate	imm	imm	32
01	Memory	(mem)	Mem[mem]	32
10	Register	\$reg	Reg[\$reg]	8
11	Register Indirect	n(\$reg)	Mem[Reg[\$reg] + n]	\$reg=8, \$n=24

The following values are used on the ISA specification:

Argument	Size	Comment
<reg>	8	Register offset.
<value>	2 + 32	Any listed addressing mode. 2 indicator bits, 32 for data.
<addr>	2 + 32	Any listed addressing mode except immediate. 2 indicator bits, 32 for data.

4 Instruction Set

Instruction	Syntax	Operation/Comments
Data Transfer		
Load	load <reg> <value>	Load a half-word (32-bit) into a register. Reg[\$reg] = \$value
Load Upper	loadu <reg> <value>	Load a half-word (32-bit) into the upper half of a register. Reg[\$reg][32:] = \$value
Load Word	loadw <reg> <value>	<i>Pseudo-instruction.</i> Loads a word (64-bit) into a register. load \$reg \$value[:32] loadu \$reg \$value[32:]
Zero	zero <reg>	<i>Pseudo-instruction.</i> Zeroes/clears a register. xor \$reg, \$reg
Store	store <reg> <addr>	Copy from register to memory. Mem[\$addr] = Reg[\$reg]
Arithmetic		
Add	add <reg> <reg> <value>	Add value to a register. Reg[\$reg1] = Reg[\$reg2] + \$value
Subtract	sub <reg> <reg> <value>	Subtract value from a register. Reg[\$reg1] = Reg[\$reg2] - \$value
Multiply	mul <reg> <reg> <value>	Multiply register by a value. Reg[\$reg1] = Reg[\$reg2] × \$value
Division	div <reg> <reg> <value>	Divide a register by a value. Reg[\$reg1] = Reg[\$reg2] ÷ \$value
Integer division	idiv <reg> <reg> <value>	Divide a register by a value, cast result to integer. Reg[\$reg1] = ⌊ Reg[\$reg2] ÷ \$value ⌋
Branching		
Compare	cmp <reg> <value>	Compare \$1 with \$2, setting comparison bits in flag register. E.g., set lt iff \$1 < \$2. Note Z flag is set depending on value, not register.
Branch	b<cond> <value>	<i>Pseudo-instruction</i> Branch to the given address if comparison matches conditional. load<cond> \$ip, \$value
Jump □	jmp <value>	<i>Pseudo-instruction.</i> load \$ip \$value

Logical		
Not	not <reg> <reg>	Bitwise NOT a register. Reg[\$reg1] = ~ Reg[\$reg2]
And	and <reg> <reg> <value>	Bitwise AND between register and value. Reg[\$reg1] = Reg[\$reg2] & \$value
Or	or <reg> <reg> <value>	Bitwise OR between register and value. Reg[\$reg1] = Reg[\$reg2] \$value
Exclusive Or	xor <reg> <reg> <value>	Bitwise exclusive-OR between register and value. Reg[\$reg1] = Reg[\$reg2] ⊕ \$value
Logical Right Shift	shr <reg> <reg> <value>	Logically shift the register right an amount. Reg[\$reg1] = Reg[\$reg2] >> \$value
Logical Left Shift	shl <reg> <reg> <value>	Logically shift the register left an amount. Reg[\$reg1] = Reg[\$reg2] << \$value
Stack		
Push	push <value>	Push a value onto the stack. Mem[Reg[sp]] = \$value Reg[sp] += 4
Push Long	pushl <value>	<i>Pseudo-instruction</i> Push a 64-bit value onto the stack. push \$value[:32] push \$value[32:]
Pop	pop [reg]	<i>Pseudo-instruction</i> Pop value from the stack, load into register if provided. sub sp, 4 If register: load \$reg, (sp)
Pop Long	popl [reg]	<i>Pseudo-instruction</i> Pop a 64-bit value from the stack, load into register if provided. sub sp, 8 If register: loadl \$reg, (sp)
Functions		
Function Call	call <value>	Call procedure at location value. More complex than load ip, \$value as pushes stack frame.
Store Arguments	stargs <value> <value> ...	<i>Pseudo-instruction</i> Push all argument values onto the stack. Useful shorthand for function call. push \$value1 ... push \$valuen push n Note assembler caches this n.
No Arguments	noargs	<i>Pseudo-instruction</i> Tells assembler that the next function call expects no arguments. push 0 Note caches n = 0.
Load Argument	ldarg <reg> i	<i>Pseudo-instruction</i> Load the i th argument (assuming all 32-bit) into the register. load \$reg, off(fp) Note see “retrieving arguments” for off calculation. Note number of arguments, n, is cached by the assembler.
Return	ret	Return from function call.
System Call	syscall <value>	Invoke the system call mapped to the given value. See the respective section for mappings.
Miscellaneous		

No-Operation \square	<code>nop</code>	Useless operation; do nothing. Equivalent to <code>or r1, 0</code> . Implemented as actual operation for efficiency.
Exit	<code>exit</code>	<i>Pseudo-instruction</i> Exit/halt the program. <code>syscall <opcode: exit></code>

4.1 Pseudo-Instructions

These are instructions which are not necessary for full functionality, but are provided for usefulness. They may be implemented using other instructions. It is up to the implementer whether to implement these as actual instructions or expand them to their equivalent form.

4.2 Instruction Layout

All instructions are encoded in a single 64-bit word. The layouts of various types is listed below. The size field stated the size in bits of this field. From top-to-bottom, the table starts at the least-significant bit.

Note, the opcode of each instruction is not decided upon; it may be any value as long as the instruction set is implemented. The only exception is `nop`, which maps to a fully-zeroed word.

Generic Layout A generic instruction layout, including the optional conditional testing. Most instructions take this format.

Bit	Purpose	Comments
0-5	Opcode	
6-9	Conditional test	These bits are tested against <code>\$flag</code> to determine if instruction is executed or skipped. <ul style="list-style-type: none"> • 1111: skip test. • 1001: test if zero flag is set. • 1000: test if zero flag is unset. Otherwise, match lower 3 bits to <code>\$flag</code> .
10-64	Instruction dependant.	

Data-Type Indicator Some instructions have a field to specify the data-type of the data being operated on. These bits are after the ordinary header, and are as follows:

Bit 0 Decimal?	Bit 1 Signed?	Bit 0 Full or half word?	Suffix	Comments
0	0	0	hu	32-bit unsigned integer.
0	0	1	[u]	64-bit unsigned integer.
0	1	0	hs	32-bit signed integer.
0	1	1	s	64-bit signed integer.
1	0	0	f	32-bit float.
1	0	1	d	64-bit double.

5 Calling Convention

Despite being a RISC processor, this processor will support explicit `call` and `ret` functions which will aid in pushing and popping a stack frame. For ease of programming, multiple actions are taken in each to maintain structure, so they are not pseudo-instructions.

5.1 Function Invocation

To call a function `[at] func` with n arguments:

```
push <arg1>
```

```
...
```

```

push <argn>
push n
call <func>

```

Note when zero arguments are needed, still **push 0** to indicate this.

Stack	
Before	After
	preserved GP registers \leftarrow \$sp
	old ip
	old fp \leftarrow \$fp
	n
	args
xxx \leftarrow \$sp	xxx

5.2 Function Returning

To return from the function invoked in the previous sub-section, we need only a call to **ret**. This will restore and pop the stack frame, as well as handle any arguments the user pushed. The following operations take place:

```

Reg[$ip] = old ip
Reg[$fp] = old fp
Reg[$sp] = loc(xxx)

```

5.3 Argument Retrieval

The frame pointer points to the top of the previous frame. Using the diagram above, it is possible to retrieve an argument from the stack. It is important to note that the size of the additional information pushed via the processor may theoretically vary, and so referencing and relying on knowledge of this size is unadvised.

i : argument index, 0-indexed; n : number of arguments.
 $\text{Arg } i = \text{Reg}[\$fp] - 4 * (1 + n - i)$

6 System Call

System calls are core functionality abstracted inside the processor. Actions are assigned operation codes and invoked via **syscall <opcode>**. Optionally, each read arguments from general-purpose registers **r1** onward.

Service	Opcode	Arguments	Operation	Result
Output				
print_int	1	\$r1 = integer	Print 64-bit integer.	None
print_float	2	\$r1 = float	Print 32-bit float.	None
print_double	3	\$r1 = double	Print 64-bit double.	None
print_char	4	\$r1 = byte	Print byte as ASCII character.	None
print_string	5	\$r1 = string address	Print null-terminated string at the address.	None
Input				
read_int	6	None	Read a signed 64-bit integer.	\$ret = integer
read_float	7	None	Read a 32-bit float.	\$ret = float
read_double	8	None	Read a 64-bit double.	\$ret = double
read_char	9	None	Read an ASCII character.	\$ret = character
read_string	10	\$r1 = string address \$r2 = max length	Read a null-terminated string into given address. String is truncated to maximum length.	None
Program Flow				
exit	11	None	Exit program.	None
exit2	12	\$r1 = exit code	Exit program with the given code.	None