

Language Documentation

Ruben Saunders

January 2025

Contents

1	Introduction	3
1.1	Overview	3
1.2	The Compiler	3
1.3	“Hello, World”	3
2	Names & Scope	4
2.1	Variables	4
2.1.1	Shadowing	4
2.1.2	Constants	4
2.2	Namespaces	5
2.2.1	Nested Namespaces	5
2.3	The Discard Symbol, ‘_’	5
2.3.1	In Let Statements	6
2.3.2	As a Parameter Name	6
3	Basic Syntax	7
3.1	Expressions & Statements	7
3.1.1	Value Categories	7
3.2	Code Blocks	7
4	Types	8
4.1	Primitive Types	8
4.2	The Unit Type	8
4.2.1	Zero-Sized Types	8
4.3	Pointer Types	9
4.4	Function Types	9
4.5	Casting	9
5	Functions	10
5.1	Calling Functions	10
5.2	Returning a Value	10
5.2.1	Unreachable Code	11

5.3	Function Declaration	11
5.4	Function Overloading	11
5.4.1	Overload Resolution	12
6	Operators	13
6.1	Casting	14
6.2	Address-Of & Dereferencing	14
6.3	Register Lookup	14
6.4	Operator Overloading	14
7	Control Flow	16
8	Modules & Imports	16
9	Standard Library	16

1 Introduction

This document will contain information about the Edel language. As an educational language, it is designed to be simple and easy to learn and operate, with the goals of being clear and readable. Each section is written from the perspective of little programming knowledge.

1.1 Overview

As the most wide-spread paradigm, this language is imperative by nature. Its syntax is inspired by C, but borrows features from C++ such as namespaces and operator overloading. However, unlike these two languages, names (except variables) are *forward declared*, meaning they can be used prior to their definition. This foregoes the need for header files of explicit forward declarations as in C and C++, and allows a much freer programming experience similar to modern languages.

1.2 The Compiler

A program consists of one file (the entry file), which is compiled by the compiler to produce an assembly output.

```
1 $ ./compiler <input_file> [-o <output_file>] [flags]
```

Where [flags] are as follows:

- `--ast` - prints the parsed program's AST.
- `--[no-]function-placeholder` - changes behaviour of a function declaration; see subsection 5.3.
- `--[no-]indentation` - adds/removes indentation from generated assembly.

Note that the output file is optional.

1.3 “Hello, World”

As is customary, a quick “hello world” program written in Edel.

```
func main() {  
    print("Hello, world\n");  
}
```

2 Names & Scope

A *name* in Edel refers to an identifier, which includes functions, variables, and namespaces. All names have a *scope*, which is the area in which that name exists/may be referenced. That is, outside its scope, a symbol does not exist.

The outermost scope is known as the *global* scope. Names defined here will be available anywhere in the current file.

The counterpart is the *local* scope, which refers to the latest scope. New scopes are created by using code blocks, for example, in while loops and functions.

2.1 Variables

Using the common analogy, a variable is like a labeled box which stores data. Each box has a *type* which tells us what it stores. For example, “int 3”. More technically, a variable is a location in memory where data can be stored, with its type determining the size of the variable and how its value can be used.

Variables are defined using the `let` keyword, with the variable’s type following a colon. For example, to define our “int 3” from earlier,

```
let foo: int = 3;
```

If required, multiple variables can be defined in the same `let` statement.

```
let foo: int = 3, bar: float = 3.14;
```

If the type is omitted, it is deduced from the assigned value. Hence, this can only be done if the variable is assigned.

```
let foo = 3, bar = 3.14;
```

(In this case, both variables adopt the types as previously written.)

2.1.1 Shadowing

Edel allows names to be *shadowed*, meaning a variable may re-use a name.

In non-local shadowing, the previous definition becomes visible again after the current scope is exited.

```
let foo = 1;
{
    let foo = 2;
    // foo = 2
}
// foo = 1
```

Edel also allows local shadowing. In this case, once shadowed, the previous definition is essentially lost and cannot be referenced.

```
let foo = 1;
// foo = 1
let foo = 2;
// foo = 2
```

2.1.2 Constants

Constants are values that, once defined, cannot be changed. Constants may be defined using the `const` keyword in place of `let`. As they cannot be re-assigned, a constant definition **must** be initialised.

```
const pi = 3.14159;
```

In the case of multiple definitions, *all* symbols will be defined as `const`.

2.2 Namespaces

Namespaces allow the programmer to group related names in an environment. Not only does this aid in organisation, but also helps to prevent naming collisions. Note that this does **not** create a new lexical scope.

A namespace is created using the `namespace` keyword, and may contain any structure which is permitted at the global level.

```
namespace maths {  
    func cos(x: double) -> double;  
}
```

To reference a namespace's member outside the namespace, it must be prefixed by the namespace's name.

```
cos(0); // error: name 'cos' does not exist  
maths.cos(0); // = 1
```

Unlike variables, namespaces cannot be shadowed. Note that the `namespace` keyword only creates a namespace if it does not already exist. If it does, then the declarations will be added to the existing namespace. For example, the below snippet would result in the definitions of the two functions `maths.sin` and `maths.cos`.

```
namespace maths {  
    func sin(x: double) -> double;  
}  
namespace maths {  
    func cos(x: double) -> double;  
}
```

2.2.1 Nested Namespaces

Namespaces may be nested explicitly as follows.

```
namespace std {  
    namespace maths {  
        func cos(x: double) -> double;  
    }  
}  
maths.cos; // error: name 'maths' does not exist  
std.maths.cos; // ok
```

However, namespace names may contain paths, that is, multiple identifiers joined by the dot `'.'` operator. This would be equivalent to explicitly nesting each identifier as a namespace. For example, one may re-write the above snippet as

```
namespace std.maths {  
    func cos(x: double) -> double;  
}
```

2.3 The Discard Symbol, `'_'`

The discard symbol, `'_'`, is a special identifier meant for instances where the result/value is intended to be discarded. It cannot be used as a general identifier, with its special uses outlined below. For example, it cannot be used as a function or namespace name.

2.3.1 In Let Statements

This symbol may be used in a let statement. In this case, an assignment is required, and the variable **cannot** have a type annotation. While the right-hand side will be evaluated, no assignment will take place.

```
let _: int; // error: expected assignment
let _: int = 1 + 2; // error: cannot have a type annotation
let _ = 1 + 2;
```

2.3.2 As a Parameter Name

Another use is for function parameters. Unlike ordinary parameters, this symbol may be re-used for multiple parameters. Then, as this symbol cannot be referenced, this is equivalent to “throwing away” the parameter.

```
func middle(_: int, x: int, _: int) -> int {
    return x;
}
```

3 Basic Syntax

Inspired by C, Edel is not whitespace sensitive, with scopes marked by braces `{}`. Semicolons are not required after every line, namely closing braces, but are required after expressions.

Comments come in two forms: single- and multi-line. Single-line comments `// ...` comment out the remainder of the current line; multi-line comments `/* ... */` comment out anything they enclose.

3.1 Expressions & Statements

Generally, a statement is some code which does not produce a value, whereas an expression does. In Edel, *everything* is an expression, with statements, by default, returning the unit type `()`.

Expressions are expected to be terminated by a semicolon. The exception to this is the final expression in a block; if omitted, the block ‘returns’ the result of this expression.

```
let a = { 1 + 2; }; // a = ()
let b = { 1 + 2 }; // b = 3
let c = { let z = 2 * b; z } + 2; // c = 8
```

Note that only the last expression in a block may have their semicolon emitted. By always terminating expressions, the block will act like a statement.

3.1.1 Value Categories

Each expression, be it an operator with operands, literals, or variable names etc., is characterised by a *type* and a *category*. Types, covered in section 4, give meaning to data and what it means in context, whereas the category tells dictates *how* a value can be used.

Edel has two categories, be them l- and r-values. The former is a value which refers to a memory location, such as variables, while the latter has no identifiable location in memory. The *l* and *r* stand for *left* and *right* because the respective values may appear on the left- or right-hand side of an assignment operator. That is,

```
a = 1; // a (lvalue) = 1 (rvalue)
```

As an rvalue has no identifiable location, it cannot be assigned to, nor can it be access via the dot `.'` operator. Instances of misuse of an l- or rvalue results in an error “expected l/rvalue, got type”.

```
1 = 2; // error: expected lvalue, got i32
true.a; // error: expected lvalue, got bool

namespace maths {}
maths + 1; // error: expected rvalue, got namespace
```

3.2 Code Blocks

As a whitespace-ignorant language, Edel uses blocks to denote structure and scope. A block is introduced by braces `{}` and contains a section of code. As previously mentioned, a block may be an expression or a statement, hence is permitted anywhere they would be expected.

Blocks in Edel introduce a new lexical scope (see section 2 for scope resolution).

4 Types

Continuing the analogy from variables, a type is a label that tells the program what kind of data a variable holds. Every variable and value has a type; a value without a type is meaningless, as a type gives data a meaning.

4.1 Primitive Types

Primitive types are a subset of those provided by the compiler representing atomic values, such as numbers.

Type	Alias	Description
<code>bool</code>		A truth value – either <code>true</code> or <code>false</code>
<code>u8</code> <code>i8</code>	<code>byte</code>	Unsigned 8-bit integer Signed 8-bit integer
<code>u16</code> <code>i16</code>		Unsigned 16-bit integer Signed 16-bit integer
<code>u32</code> <code>i32</code>	<code>int</code>	Unsigned 32-bit integer Signed 32-bit integer
<code>u64</code> <code>i64</code>	<code>long</code>	Unsigned 64-bit integer Signed 64-bit integer
<code>f32</code> <code>f64</code>	<code>float</code> <code>double</code>	32-bit floating-point number 64-bit floating-point number

Table 1: Primitive Types

These types have the following type hierarchy:

- `int8 < int16 < int32 < int64`
- `uint8 < uint16 < uint32 < uint64`
- `uintn < int(n + k)`
- `float32 < float64`
- `[u]intn < float64`
- `[u]intn < float32` where $n < 64$.

Note that `bool` is not related to anything. Unlike in other language, `bool` is treated as a distinct type with abstract true/false values, rather than a synonym for integers 1/0, respectively.

4.2 The Unit Type

The unit type, denoted by empty brackets `()`, is a special type meaning *empty*. Representing an object with no value, it is a zero-width type, meaning a symbol of this type occupies no space. For examples of use, a unit is returned from a block which returns no value.

```
let a = {}; // a: ()
```

4.2.1 Zero-Sized Types

It has been mentioned that `()` is a zero-sized type, meaning it takes up no physical space. These types exist purely for type-checking and compile-time enforcement of constraints but do not contribute to a program's memory footprint.

```
operator +(a: int, b: ()) {  
    return ();  
}  
let ans = 1 + (); // = ()
```

In the above snippet, as `a` has type `()`, it does not occupy any space on the stack. However, the assignment expression is still executed, despite the result of `()`, to ensure any side-effects are captured.

4.3 Pointer Types

The use of pointers is designed to be minimal in Edel, but they are available nonetheless. Simply put, a pointer stores the *location* (or *address*) of a variable, which can then be used to later reference and access that data. A pointer is indicated by a star `*` preceding the type; that is, type `*T` is telling the compiler *this is a location, and data of type T is stored there*.

Subsection 6.2 explains the use of the `&` and `*` unary operators to manipulate pointers.

Pointers do not fit into the type hierarchy, and cannot be implicitly cast. For safety, it is best not to cast pointer types unless absolutely necessary.

4.4 Function Types

4.5 Casting

Casting refers to the changing of a type, and can be split into two categories: implicit, and explicit. Note that the destination type of either cast **cannot** be zero-sized.

Implicit casting is where a cast is done behind-the-scenes. Type *a* may be implicitly cast (also known as *coerced*) to another type *b* if and only if $a < b$, that is, *a* is a subtype of *b*. In other words, *a* can be coerced into *b* if *b* is sufficient to represent *a* without data loss. For example, `u8` may be coerced into `u64`, or `float` into `double`, but **not** `i32` into `i16`. Implicit casting is done automatically in the following:

```
let a: double = 3.14; // float -> double
```

On the other hand, a cast which violates the type hierarchy must be *explicit*. Note that not all datatypes may be cast between each-other using explicit casting, but the subtype relationships may be largely violated.

```
let a: u8 = 42; // error!  
let a = 42 as u8; // i32 -> u8
```

5 Functions

Functions are labelled blocks of code which may be called by the programmer. A function may take some input values, known as *arguments*, and may optionally return a value.

An example of a function which takes two integers and returns an integer:

```
func add(a: int, b: int) -> int {  
    return a + b;  
}
```

A function is defined using the `func` keyword, followed by a name. Parameters in the form `name: type` are listed, comma-separated, between brackets. Note that if no arguments are required, the brackets may be omitted altogether. The function's return type is listed after an arrow; if no return type is needed, this may also be omitted. For example,

```
func foo { // signature foo() -> ()  
    // ...  
}  
  
func bar(a: float) { // signature foo(float) -> ()  
    // ...  
}
```

5.1 Calling Functions

Functions are called by following the function name with brackets `()` (this is the operator name). Any arguments to be supplied to the function are comma-separated inside these brackets, known as an *argument list*. A function with a matching signature is then found (see subsection 5.4 on function overloading), and its return value produced.

```
func add(a: int, b: int) -> int {  
    return a + b;  
}  
  
let result = add(1, 4); // result: int = 5
```

Note that `()` is an operator like any other. Therefore, if the subject is not a function type, an overload for `operator()` is searched for. The signature will be the subject type followed by each argument's type. For example,

```
5(3.14, true); // operator()(i32, f32, bool)
```

5.2 Returning a Value

The `return` keyword is used to exit/return from the current function. Optionally, the keyword may be followed by an expression, in which case that value will be returned. The type of the returned expression must match, or be a subtype of, the function's return type. Furthermore, if the function returns a non-unit type, a return statement must be present.

```
func foo {  
    return 42; // error: cannot convert i32 to ()  
}  
  
func bar -> int {  
} // error: missing return statement
```

Note that *all* code paths should return, even if you, the programmer, knows that a path is never taken. That is because the compiler cannot currently recognise, for example, always-falsy conditions.

```
func foo -> int {
    if false {
    } else {
        return 42;
    }
} // error: missing return statement
```

5.2.1 Unreachable Code

As `return` exits the function, any code after said statement is *unreachable*. This code will never be executed, and a warning will be reported.

```
func foo {
    return;
    a; // warning: unreachable code detected
}
```

The above code would usually result in an error, as symbol ‘a’ is not defined. However, as the code is not reached, no error is raised.

5.3 Function Declaration

A function declaration is one which has no body. For example,

```
func add(a: int, b: int) -> int;
```

As functions in Edel do not require forward-declaration, what does this do? Depending on the compiler flag `--function-placeholder`, it does one of two things.

Enforce Existence The default, `--no-function-placeholder`, a function declaration is used to inform the programmer about the existence of a function. One common use is in auto-generated standard library files, where the function exists but is implemented internally. If the function does not exist, an error will be raised.

```
func imaginary_function(x: int) -> bool; // error: function was declared but
does not exist
```

Generate Placeholder In this instance, if a function is declared but does not exist, the function will be created to return zero, ‘()’, false, or null, depending on the return type. This is useful when a function is yet to be implemented.

For example, the following are identical.

```
func imaginary_function(x: int) -> bool;
// ==
func imaginary_function(x: int) -> bool {
    return false;
}
```

Note that as this generates a definition, it cannot be re-defined.

5.4 Function Overloading

A function in Edel is not only identified by its name, but also its *signature*. The signature of a function is a combination of its name and parameter types. In an Edel program, a function signature must be unique, or the result would be an error. This means that multiple functions with the same name can exist. This is known as function overloading, and a variant of a function with the same name but different parameters is called an *overload*.

```
func add(a: int, b: int) -> int {  
    return a + b;  
}  
func add(a: float, b: float) -> float {  
    return a + b;  
}
```

In this case, a reference to `add` no longer just refers to one symbol. Instead, it becomes what is known as an *overload set*, and more information is required for a resolution.

```
add; // error: unable to resolve symbol 'add'
```

5.4.1 Overload Resolution

This is the process by which an overload is selected. Firstly, a list of candidates is generated; this would be all functions with a matching name. A candidate is considered a match if the number of arguments match and the type of each parameter matches. Due to subtyping rules, the latter check may lead to the generating of multiple ‘matching’ candidates. In such a case, the candidate with the highest number of identical types wins.

```
func add(a: int, b: int) -> int;  
func add(a: long, b: long) -> long;  
  
add(1, 2); // signature is add(int, int)
```

In the above snippet, both candidates match the callee. However, the former wins as `int == int` while `int < long`. If there are still ties at this point, an error is raised as a candidate cannot be decided. In this case, it is best to explicitly cast one or more arguments to reach the desired candidate.

```
func add(a: int, b: long) -> long;  
func add(a: long, b: int) -> long;  
  
add(1, 2); // error: unable to resolve symbol 'add'
```

6 Operators

Operators are special symbols which combine one or two expressions, known as unary and binary operators, respectively. In Edel, operators are generic, and any string containing the following characters is parsed as an operator:

!#\$%&*+./<=>?@^|-

As operators are generic (rather than a pre-defined set of symbols), certain information is determined by its use. Namely, an operator's *arity* is determined by its placement: unary if it is seen preceding a term, binary otherwise.

An operator's *associativity* determines the direction in which a chain of said operators are evaluated. As an example, let \sim be a binary operator in the expression $a \sim b \sim c$.

- Left-associative (ltr, \longrightarrow): equivalent to $(a \sim b) \sim c$.
- Right-associative (rtl, \longleftarrow): equivalent to $a \sim (b \sim c)$.

Finally, operators have a *precedence*. This is how tightly an operator binds, and is useful to avoid having to explicitly include brackets everywhere. For example, $a + b * c \equiv a + (b * c)$ but **not** $(a + b) * c$ as $*$ binds tighter than $+$.

The table below displays a list of built-in operators alongside their associativity and precedence. The associativity and precedence of a generic (non built-in) operator is also shown.

Table 2: Operator Precedence and Associativity

Operator	Description	Associativity	Overloadable
a.b	Member access	\longrightarrow	No
a()	Function call	\longrightarrow	Yes
(T) a &a *a registerof a	Primitive cast Address-of Dereference Register lookup	\longleftarrow	No
-a a !a <op>a	Negation Bitwise NOT Boolean NOT Generic	\longleftarrow	Yes
a as T	Cast	\longleftarrow	No
a * b a / b a % b	Multiplication Division Modulo (remainder)	\longrightarrow	Yes
a + b a - b	Addition Subtraction	\longrightarrow	Yes
a << b a >> b	Bitwise left shift Bitwise left shift	\longrightarrow	Yes
a == b a != b a < b a <= b a > b a >= b	Relational operators	\longrightarrow	Yes
a & b	Bitwise AND	\longrightarrow	Yes
a ^ b	Bitwise XOR	\longrightarrow	Yes
a b	Bitwise OR	\longrightarrow	Yes
a <op> b	Generic	\longrightarrow	Yes
a = b	Assignment	\longleftarrow	No

6.1 Casting

The concept of casting is explained in subsection 4.5. Edel offers two casting operators: `(T) a` and `a as T`.

The former is a primitive cast, meaning it only supports casting to primitive, atomic types. This cast is direct and is ignorant of the type of `a` and whether the cast is meaningful.

The latter, using the `as` keyword, is a straightforward explicit cast. There are no restrictions on `T` but, as explained in subsection 4.5, this cast still must obey some rules and will fail otherwise.

```
let a = (int) 3.14; // ok
let b = (*int) 0; // syntax error
let c = 0 as *int; // ok
```

Also note the difference in precedence: the `as` cast is slightly looser than other unary operators.

```
!1 as bool; // same as '(!1) as bool' NOT '!(1 as bool)'
```

6.2 Address-Of & Dereferencing

Both operations involve pointers, with the former generating pointers, and the latter ‘consuming’ them.

The address-of operator, `&`, calculates the memory address of a symbol. That is, if `a: T`, then `&a` is of type `*T`.

Dereferencing, then, is address-of’s twin; it retrieved the value stored at a pointer. That is, if `p: *T`, then `*a` is of type `T`.

```
let a: int = 5;
let p: *int = &a;
let b: int = *p; // then a == b
```

The dereference operator returns an lvalue, so it may be assigned to. In this case, the memory pointed to by the pointer will be written to. Continuing the example from above,

```
*p = 10; // a = 10
```

However, `b` will still be 5 as it was assigned to `*p` previously and is no longer tied to the same address.

6.3 Register Lookup

This is a lower-level operator meant for exploration of the processor’s registers.

Expecting a symbol argument, the `registerof` operator returns the index (offset) of where the given symbol is located. This allows one to gain insight into the register allocation algorithm.

For example, given the basic program below, one would expect the following.

```
let a = 5;
registerof a; // = 11, corresponding to $r1
```

If the symbol has not been loaded into registers (i.e., it has not been used at all/recently), the operator returns `-1`.

```
let a: int;
registerof a; // = -1
```

6.4 Operator Overloading

Much like functions, operators may be overridden, but this also allows new, custom operators to be defined. An operator overload is defined much like a function, except starting with the `operator` keyword, followed by the operator symbol. For example,

```
operator $(x: int) -> int {  
    return x * 100;  
}
```

defines a new unary overload for the operator \$ when applied to an integer. Similarly, one may define a binary overload.

```
operator $(a: int, b: int) -> int {  
    return $a + $b;  
}
```

Given these definitions, one could then do the following.

```
let a = $5; // = 5 * 100 = 500  
let b = 1 $ 2; // = 1 * 100 + 2 * 100 = 300
```

It is perfectly permissible to overload built-in operators. Just note that these operators have a non-standard associativity and precedence.

```
operator +(a: int, b: float) -> float {  
    return (float) a + b;  
}  
let ans = 1 + 3.14; // = 4.14
```

Like functions, an existing operator overload cannot be shadowed. For example, the following would fail.

```
operator +(a: int, b: int) -> int { // error: operator+(i32, i32) already exists  
    // ...  
}
```

- 7 Control Flow
- 8 Modules & Imports
- 9 Standard Library